
TOXIC COMMENTS CLASSIFICATION

PROJECT REPORT

Aarav Varshney and Argha Chakrabarty

Ashoka University

December 2021

Contents

1	Introduction	1
2	Literature Survey	1
3	Project Description and goals	2
4	Dataset Specification	3
5	Implementation details	3
6	Observations	5
7	Conclusions and future directions	5

1 Introduction

The project presents an application of NLP to classify texts posted on online platforms as comments as either toxic or non-toxic. While social media is a great medium for sharing opinions, it is more and more being used to bully and harass people. This project aims to classify such comments and hopefully make social media a bit more welcoming. We study the impact of tf-idf and SVMs, Decision Trees, Naive Bayes and CNNs. We evaluated our approaches on comments from the [Kaggle Toxic Comments Classification Challenge](#). Additionally we host the classifier so that users can interact with it.

2 Literature Survey

There are several papers already published on the topic of Toxic Comments Classification. This was helpful for us to understand the topic and the way experts have approached it. We have used the following papers for our analysis:

Toxic Comment Classification by Zaher et al.

The authors used Naive Bayes (as benchmark), RNN and LSTM to classify toxic comments. They trained the classifier on the same kaggle dataset.

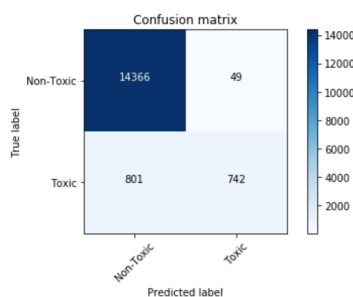


Fig. 5. Evaluation of Naive Bayes algorithm using confusion metrics.

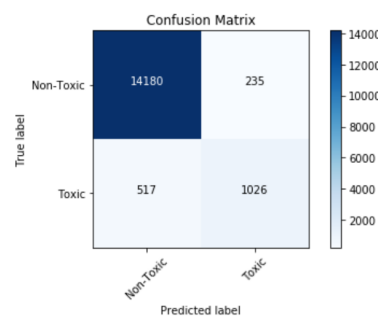
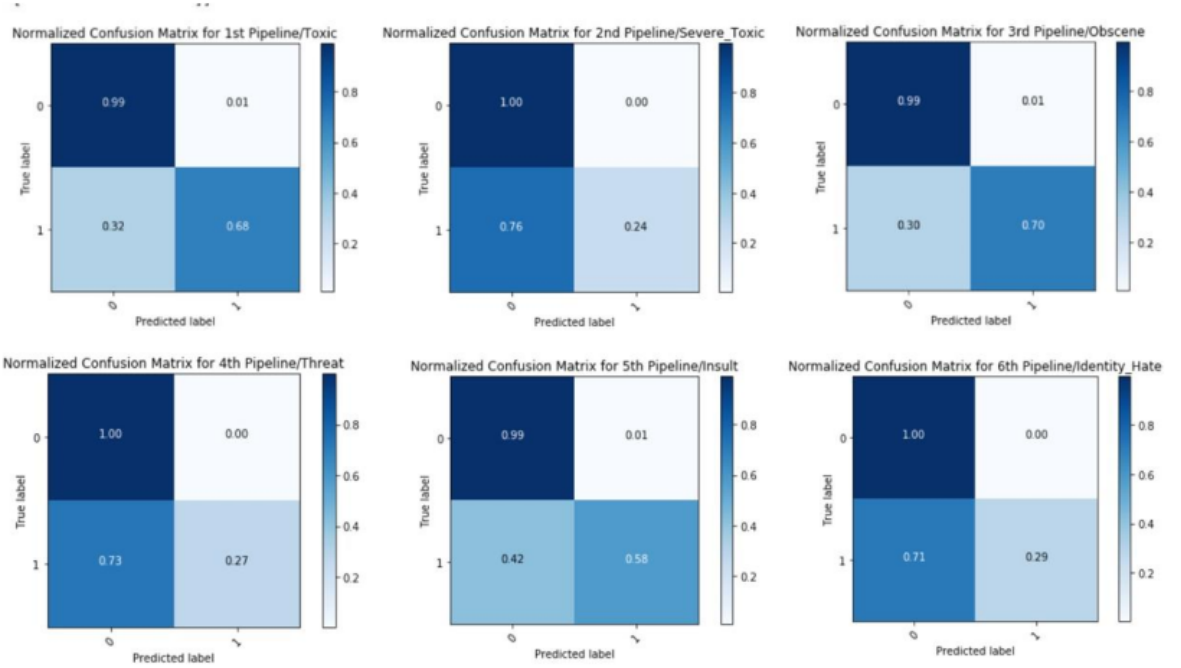


Fig. 6. Evaluation of LSTM algorithm using confusion metrics.

The results showed that LSTM has an almost 20% increase in True Positive Rate (TRP) which means that among all the identified toxic comments LSTM has 20% more sensitivity in correct assignment of the comments. Furthermore there was an increase in Recall and F1 score when using LSTM.

A ML Approach to Comment Toxicity Classification by Chakrabarty

The authors solved the problem of multi-label classification rather than binary and further classified toxic comments as hate, toxic, severelytoxic and so on. Two approaches were used for 3 labels each out of 6 labels. The authors used Bag of words using Word count vectorizer and then the tf-idf transformer which is finally used to train. Support Vector Machine (SVM) classifier with 100 as maximum iterations and C as 1.0. Whereas for the other three labels, decision tree classifier was used instead.



The paper claims that this approach performs better than LSTM and Naive Bayes and the results validates this claim. We want to validate this claim on binary classification.

3 Project Description and goals

The project is to classify toxic comments on social media platforms. The goals are as follows:

Use Different Classification Techniques

1. **Naive Bayes:** Use the results as a benchmark to compare other results with. This was chosen due to its reputation as a text classifier. This is a Generative Learning Algorithm that uses discrete values rather than continuous. Here we model $p(x|y)$ that is the probability a particular word occurs given the comment is toxic or not. We make a strong assumption that x_i 's are conditionally independent given y . Finally after fitting the model, we calculate $p(y = 1|x)$ which is just $p(x|y = 1)p(y = 1) / p(x)$
2. **Logistic Regression:** We wanted to try regression and for discrete values, we used logistic regression.
3. **Decision Trees:** Here at each level we place the best feature at the node and then the split happens depending on the features that produce the minimum classification error. And we chose the output based on the class of majority of data in the decision node.
4. **Support Vector Machines:** Support vector machine models are unique in that they find the boundaries between classes by looking at the distances between the separating line and the nearest point for each class, and are resistant to outliers.

Deploy The Model

Finally, we want to create REST API for the classifier that can be deployed on any platform. We used Flask for this.

4 Dataset Specification

The dataset was acquired from Wikipedia comments which was then labelled by Jigsaw and then finally provided by Kaggle. The dataset had 1,59,571 comments, where each comment was labelled as toxic, severe toxic, obscene, threat, insult, identity hate or none. The data was labelled by the teams at Jigsaw (Alphabet) where each comment was labelled by at least 7 humans to reduce bias. The length of each comment varied from 10 words to 1200 words. We ran an analysis to chose comments with word count less than 600 as we saw that after this length, the number of toxic comments were too few. So, a total of 1,32,327 comments were chosen.

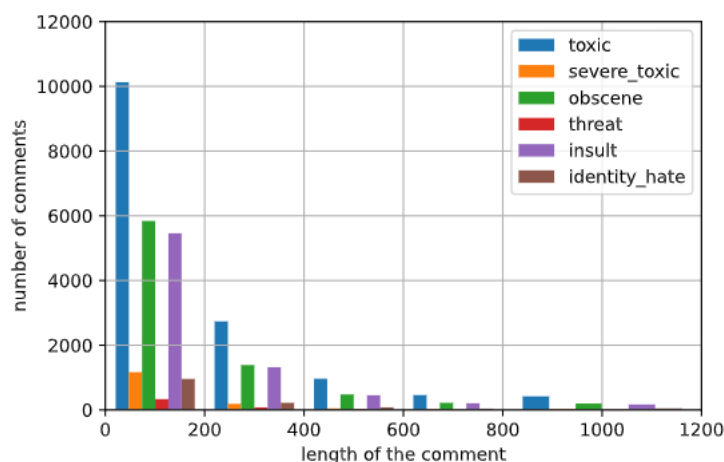


Figure 1: Length of comments

This is a multilabel classifier problem which we converted to a binary classifier problem by relabelling any comments identifying as any of the seven categories mentioned earlier. But the results for both have been published.

The dataset also went through stemming and lemmatization to remove any or all stopwords. Furthermore, we cleaned punctuations, numbers also ip-addresses as well attempted to remove usernames so that particular usernames do not make the model biased.

	id	comment_text	toxic	severe_toxic	obscene	threat	insult	identity_hate
0	0000997932d777bf	Explanation\nWhy the edits made under my usern...	0	0	0	0	0	0
1	000103f0d9cfb60f	D'aww! He matches this background colour I'm s...	0	0	0	0	0	0
2	000113f07ec002fd	Hey man, I'm really not trying to edit war. It...	0	0	0	0	0	0
3	0001b41b1c6bb37e	"\nMore\nI can't make any real suggestions on ...	0	0	0	0	0	0
4	0001d958c54c6e35	You, sir, are my hero. Any chance you remember...	0	0	0	0	0	0

Figure 2: A snapshot of the dataset

5 Implementation details

Data Preprocessing

After the data is read, we select only comments with less than 600 words in it, and label each comment as toxic or not depending on the pre-existing labels. nltk's wordnet module is used for stemming and lemmetizeing each comment after removing all the punctuations and numbers along with any other redundant utf-8 characters.

Vectorization

After all the comments are cleaned, we use TF-IDF to vectorize our comments so that we can use them for training our model. TfidfVectorizer is used as is from the scikit-learn module. TF-IDF stands for Term frequency-inverse document frequency, is a numerical statistic that is intended to reflect how important a word is to a document in a collection. A document in this case is the comment whether toxic or not. It does this by comparing the frequency of usage inside an individual comment as opposed to the entire data set. The importance increases proportionally to the number of times a word appears in the individual comment which is called Term Frequency. But if multiple comments contain the same word repeatedly (example stopwords) then their importance decreases.

$$\begin{aligned}\text{TF}(t) &= \text{Number of times term } t \text{ appears in a document} / \text{Total number of terms in the document} \\ \text{IDF}(t) &= \log_e(\text{Total number of documents} / \text{Number of documents with term } t \text{ in it}). \\ \text{Value} &= \text{TF} * \text{IDF}\end{aligned}$$

This vector becomes the feature vector for each comment. We further explored feature engineering where we can add features such as char length, word count, number of exclamations etc as features as well. We're leaving this out for now but we want to experiment more with feature engineering and see if it improves the results.

Training and Testing

The test to train split ratio is 3:1. Then we trained our model on SVM, Logistic Regression, Naive Bayes and Decision Trees. We trained the model on a laptop with i5-8520U and 16 gigs of RAM. We also trained a multi-label model using Binary Relevance from scikit-multilearn module which is built on top of scikit-learn eco-system and provides multi label classification.

Cross Validation

Cross-validation is the method used to estimate the performance (or accuracy) of machine learning models. It is used to protect against overfitting in a predictive model, particularly in a case where the amount of data may be limited. In cross-validation, you make a fixed number of folds (or partitions) of the data, run the analysis on each fold, and then average the overall error estimate. We are going with the default configuration for cross validation available in the scikit library, i.e. 5-fold cross validation.

Model	Cross Validation Score
Naive Bayes	0.920
Logistic Regression	0.953
Decision Tree	0.942
SVM	0.956

As the scores indicate, the test score of the model with cross validation is above 90% for all the models

Deployment

After training, we decided to deploy our model for demo purposes. Since we used python for our project, we decided to go with a simple FastAPI server. We created a pipeline with two stages, the first one vectorizing any input text and at the stage the actual fitting or prediction occurs. To save the model we pickle the pipeline object using python's in-built pickle library and write the raw binary data onto a file. Afterwards, the file is un-pickled at the server-side and a single string array is given to it for prediction. The API takes a single string as it's payload and returns the predicted outcome as it's response. The user has to create a POST request with the comment as the payload and we return whether the comment is toxic or not in the response body. Lastly we deployed the whole thing in a heroku server. You can take a look at the API at this link :- <https://iml-tox-det.herokuapp.com/docs> (please keep in mind that we are using a free server, it might take some time to start). We also went on and created a frontend website that uses the api: <https://adoring-bell-52f360.netlify.app>

6 Observations

Binary Classification

Model	Recall	Precision	F1-Score
Naive Bayes	0.985	0.272	0.426
Logistic Regression	0.939	0.608	0.738
Decision Tree	0.751	0.71	0.73
SVM	0.926	0.653	0.766

7 Conclusions and future directions

The results even though do not match that of the researchers, still give good results. SVM Model clearly performs better than the other models and our benchmark model. We wanted to explore some CNN models as well but since we wanted to deploy the model and use it for realtime testing, the model proved too slow for that. We also wanted to explore more feature engineering techniques and see if they would improve the results. Our future goal for the project is to test an LSTM model after fully understanding how it works and deploying it. We also want to incorporate some context detection techniques to improve the results where we analyze the context of the comment by analyzing the post, and the parent threads.

References

- [1] Zaheri, Sara; Leath, Jeff; and Stroud, David. Toxic Comment Classification, In *SMU Data Science* , Review: Vol. 3 : No. 1 , Article 13
- [2] Navoneel Chakrabarty. A Machine Learning Approach to Comment Toxicity Classification In *Jalpaiguri Government Engineering College, Jalpaiguri, West Bengal, India*
- [3] Kevin Khieu, Neha Narwal CS224N: Detecting and Classifying Toxic Comments In *Stanford University, Stanford, CA, USA*
- [4] Prinslou Tare Toxic Comment Detection and Classification In *Department of Computer Science Stanford University*